

CivicDividendOS: A Factor-Origin Fiscal Ledger and Public Automation Dividend Framework for Human–AI–Robot Economies

Author Name

Abstract—The diffusion of agentic artificial intelligence and embodied robotics is weakening the historical correspondence between human work, wages, payroll contributions, and household income. Existing proposals ask whether robots should be taxed and at what rate, or shift the fiscal burden toward capital and corporate profit; they do not resolve the prior accounting question of which factor created value when a firm or public service is jointly operated by humans, AI agents, robots, data, and conventional capital. This paper develops CivicDividendOS, a three-layer factor-origin fiscal and distribution framework. A measurement layer assigns each economically significant autonomous deployment an Autonomous Economic Activity Passport, creating an auditable economic identity—owner of record, jurisdiction, task class, resource use, and metered output—without conferring legal personality on machines. An attribution layer allocates value added across five productive factors using a Shapley-consistent estimator whose efficiency axiom guarantees budget balance and precludes double counting, and classifies each deployment as substitution, augmentation, new-task, or safety replacement. A fiscal layer applies a bounded adaptive Social Automation Contribution to the owner of record, recycles part of the proceeds through a Human Earned-Income Shield, and channels the remainder together with equity contributions into a Public Automation Wealth Fund serving citizen dividends, worker-transition insurance, human-capability investment, and public services. We state the framework’s accounting identities and establish budget balance, rate boundedness and monotonicity, a Lipschitz bound on the fiscal cost of classifier misclassification, and a closed-form condition under which the fund attains a stable ratio to output. A comparative-statics result shows that the steady-state per-capita dividend decreases in the fund’s payout ratio precisely when the fund return exceeds the growth rate of output. We then evaluate the framework in a heterogeneous-agent testbed against six policy baselines over 50 seeds. Two results qualify the design. The contribution base as first specified deducts machine costs in full against an inframarginal attributed value, which makes it marginally subsidise the automation it is meant to price; deducting at most half of those costs repairs this, and the corrected variant improves on the baseline in output, hours, labour share, inequality and the tax rate on work simultaneously, which no other arm achieves. Targeted transfers nonetheless remain superior for near-term poverty reduction, because a wealth fund accumulates over decades.

Index Terms—AI taxation, robot tax, automation, fiscal policy, public finance, value attribution, Shapley value, income distribution, universal basic capital, sovereign wealth fund, future of work, social dividend.

I. INTRODUCTION

MODERN tax systems were designed around economies in which human work generates wages, employers remit payroll contributions, households pay income and consumption taxes, and owners receive taxable profits from capital. Automation has always complicated this structure, but agentic AI and increasingly autonomous robots introduce a qualitatively broader challenge: software agents now execute cognitive, administrative, creative, analytical, and commercial tasks, while robots perform physical activities previously requiring human workers [1]–[3].

The economic output of a firm or municipal service may therefore arise from a joint production process involving humans, AI agents, robots, training data, cloud infrastructure, intellectual property, and conventional capital. If labour’s share of production falls, governments face a structural mismatch in which the tax base remains strongly linked to wages and payroll while economic value increasingly accrues to owners of automated productive capital [4], [5].

The policy question is not simply whether a robot should “pay tax.” A robot or AI agent does not possess tax residence, citizenship, consumption needs, pension rights, or legal responsibility analogous to a human worker, and current legal analysis treats AI and robots as assets, tools, services, or components of businesses rather than independent taxable persons [6]. The harder problem is how to attribute economic value produced through autonomous systems, and how to distribute a reasonable share of the resulting surplus without suppressing socially valuable innovation.

A. Three Unresolved Gaps

We identify three gaps that existing instruments do not jointly close. The *attribution gap* is that robot counts, robot investment, imputed machine wages, and firm-level profit are all poor proxies for the economic value created by autonomous systems: one AI agent may perform the work of many humans at negligible marginal cost, another may merely augment employees, a robot may perform a hazardous task that policy should encourage, and a cloud-hosted agent may serve many jurisdictions at once. The *liability gap* is that assigning machines fictional salaries and income-tax obligations is legally and administratively unworkable, because machines do not own their output and the economically relevant beneficiary is an owner, firm, platform, or operator. The *distribution gap* is that current proposals typically select a single redistribution mechanism—a robot tax, higher capital taxation, universal

basic income, retraining, or public ownership—whereas an AI-era economy plausibly requires several simultaneously, since displaced workers have immediate transition needs, future workers need education, households may need a non-wage income source, governments need stable revenue, and innovation still requires private returns.

B. Contributions

This paper makes the following contributions.

- 1) A *factor-origin accounting layer* that allocates value added across human labour, AI-agent services, robotic services, data, and conventional capital using a Shapley-consistent estimator, together with a proof that the resulting fiscal base is budget balanced and free of double counting (Proposition 1).
- 2) An *Autonomous Economic Activity Passport* that creates an auditable economic identity for machine activity and binds it to a human or corporate owner of record, thereby separating machine-level economic measurement from legal tax liability.
- 3) A *bounded adaptive Social Automation Contribution* whose rate responds to substitution, rent concentration, displacement, externalities, and base erosion, and is reduced by augmentation, training, broad-ownership, and new-task credits; we establish boundedness and monotone comparative statics (Proposition 3) and a Lipschitz bound on the fiscal cost of classifier error (Proposition 4).
- 4) A *Human Earned-Income Shield* that recycles automation revenue into reduced tax burdens on human work, with an explicit revenue-feasibility condition and a first-order behavioural correction (Proposition 5).
- 5) A *Public Automation Wealth Fund* converting part of automation rents into broadly owned capital, for which we derive the stability condition and closed-form steady state (Proposition 6) and show that the steady-state per-capita dividend is decreasing in the payout ratio precisely when the fund return exceeds output growth (Proposition 7).
- 6) A *Deployment Nexus* apportionment rule for cross-border autonomous activity that is exhaustive by construction and deliberately excludes server location as a weighting factor (Proposition 8).
- 7) A characterisation of the *marginal incidence* of the contribution base (Proposition 2), which shows that deducting machine costs in full makes the levy marginally subsidising, and a corrected base that removes the defect.
- 8) A *reproducible evaluation* in a heterogeneous-agent testbed against six baselines over 50 seeds, reported with paired bootstrap confidence intervals, together with governance safeguards on the use of simulation output in statutory rate setting.

C. Organization

Section II reviews related work and positions the framework. Section III formalizes the problem, states assumptions, and lists research questions. Section IV presents the measurement and attribution layers. Section V develops the adaptive contribution.

Section VI covers recycling and predistribution. Section VII addresses cross-border apportionment. Section VIII presents the distribution waterfall and annual fiscal cycle. Section IX describes the fiscal digital twin and its governance constraints. Section X works a numerical example end to end. Section XI specifies the evaluation protocol. Sections XII and XIII discuss implications and limitations, and Section XIV concludes.

II. RELATED WORK

A. Automation, Tasks, Wages, and Income Distribution

A central insight in the economics of automation is that technology changes the allocation of tasks between labour and capital. Acemoglu and Restrepo show that automation produces a displacement effect by allowing capital to perform tasks previously undertaken by labour, while the creation of new tasks can reinstate labour demand [1]. This task-based perspective is the natural foundation for fiscal analysis because automation affects not only aggregate productivity but also who receives the associated income, an emphasis that extends the earlier routine-task literature [7].

Empirical evidence indicates that the labour-market effects of robot adoption are uneven. Acemoglu and Restrepo document adverse employment and wage effects in U.S. local labour markets exposed to industrial robots [8], whereas Dauth et al. find that manufacturing displacement in Germany was offset by service-sector employment, with the burden falling disproportionately on younger labour-market entrants [9]. Automation can therefore generate both gains and adjustment costs, with distributional consequences that differ across time, workers, and institutional environments.

Recent AI research broadens these concerns because AI affects non-routine cognitive tasks in addition to routine physical or clerical work [2], [3], [10]. The IMF emphasizes that AI may raise productivity while widening income and wealth inequality if gains accrue disproportionately to high-income workers and capital owners [11], and OECD evidence indicates heterogeneous effects across occupations and skill groups [12].

B. Robot Taxes and Optimal Taxation

Guerreiro, Rebelo, and Teles develop a quantitative model in which falling automation costs sharply increase inequality, and find that robot taxation can be optimal during the transition while exposed workers remain in the labour force, but that the optimal tax can vanish after adjustment [13]. Thuemmel studies the joint taxation of robots, other capital, and labour income, showing that the optimal treatment depends on robot cost and general-equilibrium effects: subsidies may be desirable when robots are expensive, while a positive robot tax emerges as robots become cheaper and inequality rises [14]. Both results establish that optimal policy is dynamic rather than a permanent fixed levy.

Acemoglu et al. analyse optimal policy in a task framework and show that reducing labour taxes while combining capital-tax reductions with targeted automation taxes can improve employment outcomes relative to uniform capital taxation [15]. Related work on inefficient automation finds that firms may

automate too rapidly because they do not internalize transition costs borne by workers, creating an efficiency rationale for slowing automation even absent a redistributive objective [16]. Merola reviews taxation as a mechanism for sharing automation gains [17], and Zhang analyses automation, wage inequality, and robot taxation [18]. These contributions sit within the classical optimal-taxation tradition [19]–[21], and in particular inherit the Diamond–Mirrlees production-efficiency concern that input taxes distort factor choice unless they correct a specific externality.

The literature therefore establishes that differentiated taxation of automation may be justified under particular conditions, but a persistent implementation challenge remains: what exactly constitutes a robot, how its contribution should be measured, and how the tax can avoid discouraging beneficial capital deepening.

C. Automation, the Tax Base, and AI Public Finance

Hötte, Theodorakopoulos, and Koutroumpis examine nineteen European countries and find that the relationship between automation and tax receipts varies across technologies and diffusion phases [22]. Their results caution against assuming that automation necessarily causes a permanent decline in total revenue, while demonstrating that it can change the balance between labour, capital, and consumption tax bases. The IMF argues that fiscal policy must adapt if AI shifts income toward capital, including stronger social protection and careful consideration of capital-income taxation [4], [23]. Korinek and co-authors develop a broader public-finance framework for advanced AI economies, analysing optimal taxation under scenarios in which capital becomes increasingly dominant [5].

D. Taxing AI Rather Than Only Robots

Faivre and Cen survey the viability of AI taxation and distinguish multiple objectives, including correcting externalities, redistributing uneven gains and costs, and financing regulatory capacity; they discuss corporate income and rent-based taxation, consumption taxes on AI services, and excise-style instruments [24]. Their analysis highlights that there is unlikely to be one universal AI tax, because different harms and objectives require different bases, with measurement, incidence, international leakage, and innovation effects as the binding constraints. Legal scholarship similarly identifies substantial difficulties in treating robots or AI systems as independent taxpayers, since income-tax systems attribute income to owners, firms, or service providers rather than to the autonomous system [6]. This motivates separating machine-level economic measurement from legal tax liability, which is the organizing principle of the present framework.

E. Transfers, Basic Capital, and Public Wealth Funds

One response to automation is redistribution through universal basic income or targeted transfers [17], [25], and the ILO has examined UBI proposals against social-protection standards, highlighting financing, adequacy, labour-market, and institutional questions [26]. Transfers stabilize household

income but do not change the underlying ownership of productive capital, so a purely transfer-based model requires governments to tax a concentrated ownership base indefinitely. An alternative is broader asset ownership: Universal Basic Capital proposes granting citizens financial assets rather than only periodic income [27], and progressive sovereign wealth fund proposals argue that public ownership of capital can generate a more egalitarian distribution of investment income [28], [29]. The Alaska Permanent Fund provides a long-running example of converting resource rents into a permanent fund whose earnings support resident dividends [30]; although AI is not a natural resource, the institutional analogy is that concentrated economic rents can be partially transformed into broadly owned financial capital. Recent public-finance discussion increasingly treats equity-based mechanisms as complements to AI-era tax reform [31], [32].

F. Value Attribution and Data Dividends

Vincent et al. map design choices for data dividends and show that apparently reasonable allocation rules can generate concentrated or demographically unequal outcomes [33], and Bax investigates Shapley- and Owen-value approaches for estimating the contribution of individual data samples [34]. The cooperative-game machinery on which we draw originates with Shapley [35] and has been developed for attribution in machine learning [36]–[38]. This literature is directly useful for value attribution, but data dividends address only one input into AI production and do not provide a fiscal system for joint human–AI–robot production.

G. Adaptive and Agent-Based Fiscal Policy

AI is itself being investigated as an instrument for fiscal design. TaxAgent integrates large language models with agent-based economic simulation to adapt tax policies under heterogeneous household behaviour [39], the AI Economist applies two-level reinforcement learning to tax-schedule design [40], and other work explores AI-supported detection of tax loopholes [41]. These systems optimize conventional tax structures; they do not establish an accounting layer identifying the factor origin of value in mixed human–agent–robot production.

H. Positioning

Table I positions CivicDividendOS against the principal instrument families. The distinguishing features are that the tax base is derived from attributed value added rather than asset counts or undifferentiated profit, that the statutory taxpayer is always the owner of record, that the rate is conditioned on the measured substitution–augmentation mode, and that redistribution and predistribution operate jointly rather than as alternatives.

III. PROBLEM FORMULATION

A. Setting and Notation

Let economic output Y be jointly produced by human labour H , AI-agent services A , robotic services R , data D , and conventional capital K , so that $Y = F(H, A, R, D, K)$. Under

TABLE I
POSITIONING OF CIVICDIVIDENDOS RELATIVE TO PRINCIPAL FISCAL INSTRUMENT FAMILIES FOR AUTOMATION. “ATTRIBUTION GRANULARITY”
DENOTES THE FINEST UNIT AT WHICH VALUE CREATION IS RESOLVED.

Instrument	Tax base	Attribution granularity	Statutory payer	Sub./aug. sensitive	Pre-distribution	Cross-border rule
Fixed robot tax	Robot count / capital	Asset unit	Firm	No	No	Asset location
Capital-income tax increase	Capital income	Firm	Firm	No	No	Residence/source
Corporate rent tax	Above-normal profit	Firm	Firm	No	No	Treaty-based
UBI via general taxation	Aggregate revenue	None	Households	No	No	Domestic only
Automation tax plus retraining	Automation capital	Asset unit	Firm	Partial	No	Asset location
Public wealth fund only	Equity stakes	None	n/a	No	Yes	Ownership
Data dividend	Data contribution	Data record	Platform	No	Partial	Platform residence
CivicDividendOS	Attributed automation VA	Production episode	Owner of record	Yes	Yes	Deployment nexus

TABLE II
PRINCIPAL NOTATION.

Symbol	Meaning
$\mathcal{F}$	Factor set $\{H, A, R, D, K\}$, $n = \mathcal{F} = 5$
$p, \mathcal{P}_i$	Production episode; episodes of entity i
V_p	Value added of episode p
$v_p(S)$	Coalition value of factor subset $S \subseteq \mathcal{F}$
$\psi_{p,f}$	Shapley value of factor f in episode p
$\phi_{p,f}$	Normalized attribution share, $\psi_{p,f}/V_p$
Z_p	Automation mode of episode p
S_i	Substitution intensity index of entity i
A_i^{aug}	Augmentation benefit index
ACB_i	Automation Contribution Base
r_i	Social Automation Contribution rate
r_0	Base rate; $r_{\min}, r_{\max}$ rate floor and ceiling
SAC_i	Social Automation Contribution liability
κ	Share of SAC revenue recycled to the shield
ω	Share of SAC revenue paid into the fund
F_t	Fund value at t ; $f_t = F_t/Y_t$
r^f	Real fund return; g real output growth
ρ	Fund payout ratio (share of gross return)
s	Fund inflow as a share of output
N_t	Eligible resident population
Θ	Policy parameter vector

a traditional fiscal system, a substantial share of public revenue is generated from

$$\text{Tax}^{\text{trad}} = \tau_L W_H + \tau_K \Pi + \tau_C C, \quad (1)$$

where W_H is human wage income, Π taxable capital or corporate income, and C consumption. As autonomous factors replace or augment human tasks, the wage share may decline while output is constant or rising,

$$\frac{W_H}{Y} \downarrow, \quad \frac{Y_A + Y_R}{Y} \uparrow, \quad (2)$$

so that if ownership of A and R is concentrated, household purchasing power, payroll-financed social protection, and government revenue can all weaken even as aggregate productivity rises. Table II summarizes the notation used throughout.

B. Assumptions

The framework rests on the following assumptions, each of which is revisited in Section XIII.

Assumption 1 (Registrability). Every economically significant autonomous deployment can be registered, and its outputs

metered, at the level of a production episode. Deployments below a de minimis threshold are exempt.

Assumption 2 (Observable value added). Episode-level value added V_p is observable from accounting records, and a coalition value function $v_p : 2^{\mathcal{F}} \rightarrow \mathbb{R}_{\geq 0}$ with $v_p(\emptyset) = 0$ and $v_p(\mathcal{F}) = V_p$ can be estimated.

Assumption 3 (Monotone production). v_p is monotone, i.e. $v_p(S) \leq v_p(S \cup \{f\})$ for all $S \subseteq \mathcal{F}$ and $f \notin S$: adding a factor never reduces attainable value.

Assumption 4 (Owner of record). Every registered deployment has a unique legally responsible owner, operator, or beneficiary. No legal personality is conferred on machines.

Assumption 5 (Bounded classification error). The substitution–augmentation classifier misclassifies with probability at most ϵ , and index inputs are normalized to $[0, 1]$.

Assumption 6 (Policy horizon). Over the rate-adjustment horizon, firms are price takers in factor markets and the fund’s real return r^f and output growth g are constant. This assumption is used only for the steady-state results of Section VI.

C. Research Problem and Questions

The fundamental research problem is how a tax-and-distribution system can identify the economic value created by humans, AI agents, robots, data, and conventional capital; assign legal fiscal responsibility to the appropriate human or corporate beneficiaries; and distribute part of the automation-generated surplus in a way that preserves innovation, public revenue, household income, workforce transition, and broad participation in future capital returns. The problem is harder when AI agents are cloud-hosted, operate across borders, change tasks rapidly, and act primarily as complements rather than substitutes.

This decomposes into eight research questions. **RQ1**: can task-level factor-origin accounting estimate attributable value more accurately than robot counts, imputed robot wages, or capital expenditure? **RQ2**: can a substitution–augmentation classifier distinguish automation that displaces taxable labour from automation that complements workers? **RQ3**: does a dynamic Social Automation Contribution preserve government revenue with less distortion to innovation than a fixed robot tax? **RQ4**: can combining automation contributions with a

Human Earned-Income Shield improve employment and after-tax wage outcomes relative to raising labour or consumption taxes? **RQ5**: does a Public Automation Wealth Fund create a more durable and equitable distribution of AI-era wealth than transfers financed entirely from annual revenue? **RQ6**: what allocation among citizen dividends, worker transition, human-capability investment, and public services maximizes long-run social welfare? **RQ7**: how should cross-border AI-agent activity be apportioned when ownership, compute, users, and deployment occur in different jurisdictions? **RQ8**: how robust is the framework to avoidance, misclassification, strategic task decomposition, offshoring, and rapid capability change?

IV. THE CIVICDIVIDENDOS FRAMEWORK

CivicDividendOS is organized into a measurement layer, an attribution layer, and a fiscal layer, with a governance and simulation loop closing over the policy vector. Fig. 1 shows the architecture, and Table III maps each component to the gap it addresses and the observable it requires.

A. Measurement Layer: Autonomous Economic Activity Passport

Every economically significant autonomous deployment j is assigned an Autonomous Economic Activity Passport,

$$\text{AEAP}_j = \{\text{ID}_j, O_j, J_j, T_j, E_j, Q_j, X_j, M_j\}, \quad (3)$$

where O_j is the legally responsible owner or operator, J_j the deployment jurisdiction, T_j the task category, E_j energy and resource consumption, Q_j metered output, X_j the risk classification, and M_j the model or robot class. The record (3) is not legal personhood; it is a fiscal measurement record that allows economic activity generated by an AI agent or robot to be audited without treating the system as a taxpayer. This construction is what operationalizes the separation of measurement from liability identified as necessary in [6].

B. Attribution Layer: Factor-Origin Ledger

For each production episode p , the ledger estimates the causal contribution of the five factors. Let $v_p(S)$ denote the value attainable by coalition $S \subseteq \mathcal{F}$ under Assumption 2. The attribution to factor f is the Shapley value [35]

$$\psi_{p,f} = \sum_{S \subseteq \mathcal{F} \setminus \{f\}} \frac{|S|!(n - |S| - 1)!}{n!} [v_p(S \cup \{f\}) - v_p(S)], \quad (4)$$

with normalized shares $\phi_{p,f} = \psi_{p,f}/V_p$ for $V_p > 0$ and attributed value $V_{p,f} = \phi_{p,f}V_p$. Coalition values may be estimated through controlled counterfactuals, production-function estimation, matched historical baselines, or digital-twin simulation, following the estimation strategies developed for attribution in machine learning [36]–[38].

Proposition 1 (Budget balance and no double counting). *Under Assumptions 2 and 3, the attribution defined by (4) satisfies $\sum_{f \in \mathcal{F}} V_{p,f} = V_p$ and $V_{p,f} \geq 0$ for every f . Consequently the*

aggregate automation base satisfies $VA_i^{A,R} \leq \sum_{p \in \mathcal{P}_i} V_p$, so no unit of value added can be assessed more than once.

Proof. The Shapley value satisfies the efficiency axiom, $\sum_f \psi_{p,f} = v_p(\mathcal{F}) - v_p(\emptyset) = V_p$, which gives $\sum_f \phi_{p,f} = 1$ and hence the first identity. Under Assumption 3 every marginal contribution $v_p(S \cup \{f\}) - v_p(S)$ is non-negative, and $\psi_{p,f}$ is a convex combination of such terms, so $\psi_{p,f} \geq 0$. Summing the non-negative terms $V_{p,A} + V_{p,R} \leq V_p$ over $p \in \mathcal{P}_i$ yields the bound. $\square$

Remark 1 (Computational cost). With $n = 5$ factors, exact evaluation of (4) requires only $2^n - 1 = 31$ coalition values per episode, so no sampling approximation is needed at the factor level. Approximation is required only for the second-stage allocation within a factor when many distinct agents or robots must be separated, where standard Shapley samplers apply [38].

This is the central departure from conventional robot taxation: the system taxes neither the physical count of robots nor a fictional machine wage, but the estimated economic origin of value.

C. Substitution–Augmentation Classification

Automation that replaces a worker and automation that makes a worker more productive should not receive identical fiscal treatment. Each episode is assigned a mode

$$Z_p \in \{\text{SUBSTITUTION, AUGMENTATION, NEWTASK, SAFETYREPLACEMENT}\}, \quad (5)$$

with a substitution score $S_p = \Pr(\Delta H_p < 0 \mid A_p, R_p, \mathbf{X}_p)$ estimated from employment, hours, and task panels, and a complementary augmentation score measuring gains to human productivity or wages. Entity-level indices S_i and A_i^{aug} are value-weighted averages of episode scores taken under the mode assignment (5). This classification enters the rate directly, and its error properties are analysed in Proposition 4.

V. ADAPTIVE SOCIAL AUTOMATION CONTRIBUTION

A. Contribution Base and Statutory Taxpayer

The Automation Contribution Base of entity i is

$$\text{ACB}_i = VA_i^{A,R} - C_i^{\text{all}}, \quad VA_i^{A,R} = \sum_{p \in \mathcal{P}_i} (V_{p,A} + V_{p,R}), \quad (6)$$

where C_i^{all} comprises verified operating, compute, energy, depreciation, and acquisition costs admitted by policy. By Assumption 4 the statutory taxpayer is Owner(AEAP _{i}), never the AI agent or robot.

B. Marginal Incidence of the Contribution Base

The base (6) deducts the full cost of machine services from an attributed value that is inframarginal, and these two quantities do not have the same derivative. This has a consequence that is easy to miss and that determines whether the instrument discourages automation at all.




Fig. 1. CivicDividendOS architecture. Solid arrows carry measured or assessed quantities; dashed arrows carry policy parameters and simulation feedback. The owner of record is the sole statutory taxpayer; no fiscal obligation attaches to a machine. The digital twin proposes but does not autonomously set the policy vector Θ (Section IX).

TABLE III
FRAMEWORK COMPONENTS, THE GAP EACH ADDRESSES, AND THE OBSERVABLE EACH REQUIRES.

Component	Gap addressed	Mechanism	Required observable
Autonomous Economic Activity Passport	Liability	Registry binding activity to owner of record	Registration, jurisdiction, task class, resource use
Factor-Origin Ledger	Attribution	Shapley allocation over five factors	Episode value added, coalition counterfactuals
Substitution–Augmentation Classifier	Attribution	Mode assignment Z_p and index construction	Employment, hours, wage and task-level panels
Automation Contribution Base	Attribution	Attributed value added net of allowable cost	Verified compute, energy, depreciation costs
Social Automation Contribution	Distribution	Bounded adaptive rate on the base	Substitution, rent, displacement, credit indices
Human Earned-Income Shield	Distribution	Recycling to labour-income tax relief	Payroll base, labour supply elasticity
Public Automation Wealth Fund	Distribution	Fund accumulation and payout rule	Fund return, inflow share, population
Automation Ownership Units	Distribution	Equity contribution against public support	Subsidy, compute, data, procurement receipts
Deployment Nexus	Attribution	Multi-factor jurisdictional apportionment	Usage, revenue, users, physical operations
Fiscal Digital Twin	Distribution	Ex ante screening of policy vectors	Calibrated heterogeneous-agent model

Proposition 2 (Marginal incidence of the contribution base). *Let attribution be the Shapley value over $\{L, M\}$ of the production game with $v(\{L, M\}) = Y(L, M)$, and let the firm be a price taker choosing M so that $\partial Y/\partial M = p_M$. For the base $ACB = \psi_M - \varphi p_M M$ with deduction share $\varphi \in [0, 1]$,*

$$\frac{\partial ACB}{\partial M} = \frac{1}{2} \frac{\partial v(\{M\})}{\partial M} + \left(\frac{1}{2} - \varphi\right) p_M. \quad (7)$$

Under full deduction $\varphi = 1$ this equals $\frac{1}{2}[\partial v(\{M\})/\partial M - p_M]$, which is negative whenever the machine's stand-alone marginal product falls below its price, so the contribution is marginally subsidising rather than corrective. Any $\varphi \leq 1/2$ restores $\partial ACB/\partial M \geq 0$.

Proof. Since $\psi_M = \frac{1}{2}v(\{M\}) + \frac{1}{2}[v(\{L, M\}) - v(\{L\})]$ and $v(\{L\})$ does not depend on M , differentiating gives $\partial \psi_M/\partial M = \frac{1}{2}\partial v(\{M\})/\partial M + \frac{1}{2}\partial Y/\partial M$. Substituting the first-order condition $\partial Y/\partial M = p_M$ and subtracting φp_M yields (7). Setting $\varphi = 1$ and $\varphi = 1/2$ gives the two stated cases. $\square$

Proposition 2 is a design warning rather than a curiosity. The Shapley term ψ_M credits the machine factor with half of the complementarity surplus it shares with labour, which is inframarginal, whereas the cost deduction $p_M M$ scales one-for-one with the marginal decision. Deducting costs in full therefore makes an additional machine reduce the taxable base, and a levy on that base lowers the effective machine price. The condition under which this bites is structural rather than empirical, and it is worth stating precisely. Under the CES-with-span-of-control technology of Section XI-B the ratio of the stand-alone to the joint marginal product is exactly $s_M^{\gamma/\rho-1}$ with $s_M \in (0, 1)$ and $\rho = 1 - 1/\sigma$, so $\partial v(\{M\})/\partial M < p_M$ holds *identically* whenever $\gamma > 1 - 1/\sigma$. At the calibration of Section XI-B ($\gamma = 0.85$, $\sigma = 1.5$) the inequality therefore cannot fail, and its universal incidence across firms and quarters is a property of the calibration rather than evidence about the world. We report it as such. The implied effective wedge at that calibration is between -3.3% and -4.1% of the machine price. We therefore evaluate both the base exactly as specified in (6) and a corrected base with $\varphi = 1/2$, and recommend the

latter.

C. Rate Function

Let the raw rate score be

$$\tilde{r}_i = r_0 + \alpha S_i + \beta C_i^{\text{rent}} + \gamma E_i^{\text{disp}} + \delta X_i^{\text{ext}} + \eta R_i^{\text{rev}} - \theta A_i^{\text{aug}} - \lambda T_i^{\text{train}} - \mu B_i^{\text{broad}} - \nu N_i^{\text{new}}, \quad (8)$$

with all indices normalized to $[0, 1]$ and all weights non-negative, where S_i is labour substitution intensity, C_i^{rent} an economic-rent and concentration score, E_i^{disp} measured worker-displacement burden, X_i^{ext} environmental or social externalities, R_i^{rev} erosion of the payroll and social-contribution base in the sense of (2), A_i^{aug} human augmentation benefit, T_i^{train} verified worker training and transition investment, B_i^{broad} a broad-ownership and benefit score, and N_i^{new} creation of new human tasks. The applied rate and liability are

$$r_i = \text{clip}(\tilde{r}_i, r_{\min}, r_{\max}), \quad \text{SAC}_i = r_i \cdot \text{ACB}_i. \quad (9)$$

A differentiable alternative, $r_i = r_{\min} + (r_{\max} - r_{\min}) \sigma(\tilde{r}_i)$ with σ logistic, may be preferred when the policy vector is tuned by gradient-based search; all results below hold for either map, since both are non-decreasing and Lipschitz with constant at most one after scaling.

Proposition 3 (Boundedness and monotone comparative statics). *For any index vector, $r_i \in [r_{\min}, r_{\max}]$ and hence the effective burden on attributed automation value added never exceeds $r_{\max}$. Moreover r_i is non-decreasing in S_i , C_i^{rent} , E_i^{disp} , X_i^{ext} and R_i^{rev} , and non-increasing in A_i^{aug} , T_i^{train} , B_i^{broad} and N_i^{new} ; the inequalities are strict wherever $\tilde{r}_i \in (r_{\min}, r_{\max})$.*

Proof. $\text{clip}(\cdot, r_{\min}, r_{\max})$ maps $\mathbb{R}$ onto $[r_{\min}, r_{\max}]$ and is non-decreasing, and $\tilde{r}_i$ in (8) is affine with non-negative coefficients on the first group and non-positive coefficients on the second. A non-decreasing map composed with an affine function preserves the sign of each partial derivative, and on the open interval where the clip is inactive the composition is the identity, giving strictness. Since $\text{SAC}_i = r_i \text{ACB}_i$ and $\text{ACB}_i \leq V A_i^{\text{A,R}}$, the burden ratio $\text{SAC}_i / V A_i^{\text{A,R}} \leq r_{\max}$. $\square$

Proposition 3 supplies an investment-relevant guarantee: whatever the index realizations, the marginal fiscal burden on automation value added is capped at $r_{\max}$, which can be published ex ante. A complementary design condition ensures that purely complementary automation is untaxed at the floor: if $\theta \geq r_0 + \beta C_i^{\text{rent}} + \eta R_i^{\text{rev}} - r_{\min}$ whenever $S_i = 0$ and $A_i^{\text{aug}} = 1$, then augmentation-only deployments face $r_i = r_{\min}$.

Fig. 2 plots the applied rate against substitution intensity for three augmentation levels under the calibration of Section X, illustrating both the monotonicity of Proposition 3 and the activation of the floor at high augmentation.




Fig. 2. Applied Social Automation Contribution rate as a function of substitution intensity, for three augmentation levels, under the calibration of Table V. Curves are evaluated directly from (8)–(9) with the displacement index tracking substitution as $E_i^{\text{disp}} = 0.10 + 0.60 S_i$; the floor $r_{\min} = 0$ binds at low substitution when augmentation is high.

D. Robustness to Classifier Error

Because the rate depends on an estimated classification, the fiscal consequence of misclassification must be bounded.

Proposition 4 (Lipschitz bound on misclassification cost). *Under Assumption 5, for any two index vectors differing only in the substitution and augmentation coordinates, $|r_i - r'_i| \leq \alpha |S_i - S'_i| + \theta |A_i^{\text{aug}} - A_i^{\text{aug}'}|$. Consequently the expected rate error induced by classification error at rate ε satisfies $|\mathbb{E}[r_i] - r_i^*| \leq \varepsilon(\alpha + \theta)$, and the expected revenue error satisfies $|\mathbb{E}[\text{SAC}_i] - \text{SAC}_i^*| \leq \varepsilon(\alpha + \theta) \text{ACB}_i$.*

Proof. The clip map is non-expansive, $|\text{clip}(x) - \text{clip}(y)| \leq |x - y|$. Applying this to (8) and using the triangle inequality on the two differing affine terms gives the first bound. Since indices lie in $[0, 1]$, each coordinate difference is at most one, so a misclassified case incurs a rate error of at most $\alpha + \theta$; averaging over a misclassification event of probability at most ε and a correctly classified event of zero error yields the stated expectation bound, and multiplying by the base gives the revenue bound. $\square$

Under the calibration of Section X, where $\alpha = 0.10$ and $\theta = 0.09$, Proposition 4 bounds the expected rate error at 0.95, 1.90, and 3.80 percentage points for classifier error rates of 5%, 10%, and 20% respectively. This is a design-relevant result: the weights α and θ can be selected jointly with a target on the tolerable fiscal consequence of classification error, so that a noisy classifier bounds rather than propagates uncertainty.

VI. RECYCLING AND PREDISTRIBUTION

A. Human Earned-Income Shield

A core design feature is that part of automation revenue is used to shrink the labour term $\tau_L W_H$ of (1) rather than to expand it. Let the labour-income tax relief satisfy

$$\Delta \tau_L(t) = -\kappa \frac{\text{RevSAC}(t)}{B_L(t)}, \quad (10)$$

where $B_L(t)$ is the labour tax base and $\kappa \in [0, 1]$ the recycling share, so that qualifying low- and middle-income

workers face lower payroll or earned-income burdens as autonomous production contributes more to public revenue.

Proposition 5 (Feasibility and behavioural correction). *The joint allocation of SAC revenue to the fund and the shield is feasible if and only if $\omega + \kappa \leq 1$. Ignoring behavioural response, the mechanical revenue cost of the shield is $\kappa \text{Rev}_{\text{SAC}}$. To first order in the labour supply elasticity ε_L with respect to the net-of-tax rate, the net cost is*

$$\kappa \text{Rev}_{\text{SAC}} \left[1 - \varepsilon_L \frac{\tau_L}{1 - \tau_L} \right]. \quad (11)$$

Proof. Feasibility is the accounting constraint that allocated shares cannot exceed collected revenue. For the correction, the shield raises the net-of-tax rate by $\Delta(1 - \tau_L) = \kappa \text{Rev}_{\text{SAC}} / B_L$, so the induced base expansion is $\Delta B_L = \varepsilon_L B_L \Delta(1 - \tau_L) / (1 - \tau_L) = \varepsilon_L \kappa \text{Rev}_{\text{SAC}} / (1 - \tau_L)$, which is taxed at τ_L and yields a revenue offset of $\varepsilon_L \kappa \text{Rev}_{\text{SAC}} \tau_L / (1 - \tau_L)$. Subtracting from the mechanical cost gives (11). $\square$

Equation (11) makes the self-financing threshold explicit: the shield is fully self-financing when $\varepsilon_L \geq (1 - \tau_L) / \tau_L$, which for $\tau_L = 0.40$ requires $\varepsilon_L \geq 1.5$, well above conventional estimates. The shield should therefore be presented as a redistribution of the fiscal burden across factors, not as a costless reform.

B. Public Automation Wealth Fund

A fixed fraction of SAC revenue and selected equity contributions enter a professionally managed Public Automation Wealth Fund, whose value evolves as

$$F_{t+1} = F_t(1 + r^f) + \omega \text{Rev}_{\text{SAC},t} + E_t - \text{Dist}_t, \quad (12)$$

with an endowment-style payout rule $\text{Dist}_t = \rho r^f F_t$ distributing a share ρ of the gross return. Writing $f_t = F_t / Y_t$ and letting the combined inflow satisfy $(\omega \text{Rev}_{\text{SAC},t} + E_t) / Y_t = s$, we obtain the following.

Proposition 6 (Fund stability and steady state). *Under Assumption 6, the fund-to-output ratio follows $f_{t+1} = af_t + s$ with $a = [1 + (1 - \rho)r^f] / (1 + g)$. A finite steady state exists if and only if*

$$(1 - \rho)r^f < g, \quad (13)$$

in which case

$$f^* = \frac{s(1 + g)}{g - (1 - \rho)r^f}, \quad (14)$$

convergence is geometric at rate a , and the steady-state per-capita dividend expressed as a share of per-capita output is $d^* = \rho r^f f^*$.

Proof. Dividing (12) with $\text{Dist}_t = \rho r^f F_t$ by $Y_{t+1} = (1 + g)Y_t$ gives the stated affine recursion. It converges if and only if $|a| < 1$, which for positive parameters is $1 + (1 - \rho)r^f < 1 + g$, i.e. (13). The fixed point of $f = af + s$ is $f^* = s / (1 - a)$, and substituting $1 - a = [g - (1 - \rho)r^f] / (1 + g)$ yields (14). The dividend follows from $\text{Dist}_t / N_t = \rho r^f F_t / N_t$ divided by Y_t / N_t . $\square$




Fig. 3. Steady-state fund-to-output ratio f^* from (14) against the payout ratio ρ , for three real fund returns, with $g = 0.03$ and inflow share $s = 0.0065$. Each curve diverges as ρ approaches the stability boundary $\rho = 1 - g/r^f$ of (13).

Proposition 7 (Payout ratio and the return–growth gap). *Under the conditions of Proposition 6, the steady-state per-capita dividend share satisfies*

$$d^*(\rho) = \frac{\rho r^f s(1 + g)}{(g - r^f) + \rho r^f}, \quad \frac{\partial d^*}{\partial \rho} = \frac{r^f s(1 + g)(g - r^f)}{[(g - r^f) + \rho r^f]^2}. \quad (15)$$

Hence d^* is strictly decreasing in ρ if $r^f > g$, strictly increasing if $r^f < g$, and independent of ρ if $r^f = g$.

Proof. Substituting (14) into $d^* = \rho r^f f^*$ and writing the denominator as $g - (1 - \rho)r^f = (g - r^f) + \rho r^f$ gives (15). Differentiating the quotient $A\rho / (B + \rho r^f)$ with $A = r^f s(1 + g)$ and $B = g - r^f$ yields $AB / (B + \rho r^f)^2$, whose sign is that of $B = g - r^f$. $\square$

Proposition 7 is the framework’s principal counterintuitive result and has a direct policy reading. In the empirically relevant regime in which the fund’s real return exceeds output growth [29], raising the payout ratio increases dividends only transitionally and lowers them permanently, because the higher payout erodes the capital base faster than it augments the flow. A fund designed for durable dividends must therefore retain a substantial share of its return, and the political pressure to raise ρ is precisely the pressure that undermines the long-run dividend. Fig. 3 shows the steady-state fund ratio as a function of the payout ratio for three return assumptions, and Table IV reports the corresponding dividends and convergence half-lives.

Inverting (14) gives a direct design rule: to sustain a dividend equal to a share κ_d of per-capita output, the required inflow is $s = \kappa_d [g - (1 - \rho)r^f] / [\rho r^f (1 + g)]$. For $\kappa_d = 0.02$, $r^f = 0.05$, $\rho = 0.60$ and $g = 0.03$, this requires a steady-state fund of 66.7% of output financed by an annual inflow of 0.65% of output. Table IV also records a finding that tempers the proposal: convergence is slow, with a half-life of roughly 71 years under this calibration. Dividend adequacy in the first decades therefore depends on the direct-transfer channel rather than on fund returns, and the fund should be presented as an intergenerational instrument rather than a short-run income policy.

TABLE IV

STEADY-STATE FUND RATIO f^* , PER-CAPITA DIVIDEND d^* AS A PERCENTAGE OF PER-CAPITA OUTPUT, AND CONVERGENCE HALF-LIFE, COMPUTED FROM PROPOSITIONS 6–7 WITH $g = 0.03$ AND $s = 0.0065$. VALUES ARE ANALYTICALLY DERIVED, NOT SIMULATED. NOTE THAT d^* FALLS AS ρ RISES, CONSISTENT WITH PROPOSITION 7 SINCE $r^f > g$ THROUGHOUT.

r^f	ρ	f^*	d^* (%)	Half-life (years)
0.04	0.55	0.558	1.23	59.1
0.04	0.75	0.335	1.00	35.3
0.04	0.95	0.239	0.91	25.1
0.05	0.55	0.893	2.46	94.8
0.05	0.75	0.383	1.44	40.4
0.05	0.95	0.244	1.16	25.6
0.06	0.55	2.232	7.36	237.6
0.06	0.75	0.446	2.01	47.2
0.06	0.95	0.248	1.41	26.1

C. Automation Ownership Units

CivicDividendOS adds a predistribution mechanism rather than relying exclusively on ex post transfers. When a qualifying firm receives major public AI subsidies, public compute, exclusive public data access, or government procurement support, policy may require a small Automation Ownership Unit contribution $\text{AOU}_i = \xi_i \cdot \text{Equity}_i$, with the units entering the fund rather than individual government departments. The logic parallels a resource fund: the public contributes infrastructure, research, data, procurement demand, education, and legal institutions that help create private automation value, and may therefore hold a modest diversified claim on future returns. Because equity contributions are non-cash, they relax the inflow constraint in (14) without increasing the current-year burden on firm liquidity, which is a distinct advantage over cash levies for capital-constrained entrants.

VII. CROSS-BORDER DEPLOYMENT NEXUS

Cloud AI complicates conventional tax residence because an agent may be trained in one jurisdiction, owned in a second, hosted in a third, and economically deployed in a fourth. We therefore define, for entity i and jurisdiction j ,

$$\text{Nexus}_{ij} = w_1 U_{ij} + w_2 R_{ij} + w_3 P_{ij} + w_4 O_{ij}, \quad (16)$$

where U_{ij} , R_{ij} , P_{ij} and O_{ij} are jurisdiction j 's shares of entity i 's metered usage, revenue, users, and physical operations respectively, and $\sum_k w_k = 1$ with $w_k \geq 0$.

Proposition 8 (Exhaustive apportionment). *If each component is a share satisfying $\sum_j U_{ij} = \sum_j R_{ij} = \sum_j P_{ij} = \sum_j O_{ij} = 1$, then $\sum_j \text{Nexus}_{ij} = 1$, and the apportioned liabilities $\text{SAC}_{ij} = \text{Nexus}_{ij} \text{SAC}_i$ satisfy $\sum_j \text{SAC}_{ij} = \text{SAC}_i$ exactly.*

Proof. Summing (16) over j gives $\sum_k w_k \cdot 1 = 1$; multiplying by SAC_i gives the second identity. $\square$

Server location is deliberately excluded from (16), since it is the most cheaply relocated of the candidate factors and therefore the most exposed to profit shifting. Revenue apportionment

can then follow verified economic deployment rather than infrastructure siting. We stress that this mechanism would require coordination with existing international corporate tax rules, in particular the OECD/G20 two-pillar framework [42], and is proposed as a research mechanism rather than a claim about current law. Robustness to strategic manipulation of the remaining four components is an open question that we do not resolve here.

VIII. DISTRIBUTION WATERFALL AND ANNUAL FISCAL CYCLE

A. Waterfall

The complete distribution mechanism begins from the factor decomposition $Y = Y_H + Y_A + Y_R + Y_D + Y_K$, with private primary distribution $Y^{\text{priv}} = W_H + \Pi_O + R_D + R_K$ unchanged. CivicDividendOS then introduces a secondary and ownership-based distribution in which households receive

$$Y_H^{\text{fin}} = W_H - \text{Tax}_H + \text{Shield}_H + \text{Dividend}_H + \text{Transition}_H + \text{Services}_H, \quad (17)$$

while owners and the public sector respectively receive

$$Y_O^{\text{fin}} = \Pi_O - \text{CIT}_O - \text{SAC}_O + \text{PrivRet}_O, \\ Y_P = \text{CIT} + \text{SAC} + \text{FundRet}, \quad (18)$$

with Y_P in (18) split among public services, human-capability investment, transition support, and citizen dividends. Annual distributable returns divide as $\text{Dist}_t = D_t^{\text{cit}} + D_t^{\text{tr}} + D_t^{\text{cap}} + D_t^{\text{pub}}$, with the per-capita civic dividend $\text{CADiv}_t = \rho D_t^{\text{cit}} / N_t$ and transition support $\text{WTI}_i = g(\text{WageLoss}_i, \text{Tenure}_i, \text{Exposure}_i, \text{Training}_i)$. The design preserves private profit incentives while creating a parallel public claim on part of the automation surplus. Fig. 4 traces a unit of value added through the full sequence.

B. Annual Fiscal Cycle

Algorithm 1 states the operational cycle. Steps 2–4 are the measurement and attribution layers, steps 5–8 the assessment layer, and steps 9–12 distribution and governance. The cycle is deliberately annual: the digital twin proposes a revised policy vector, but adoption requires the publication and oversight step described in Section IX.

IX. FISCAL DIGITAL TWIN AND GOVERNANCE SAFEGUARDS

Before changing rates or distribution shares, policy is screened in a fiscal digital twin containing heterogeneous households, firms, AI agents, robots, sectors, and jurisdictions, following the tradition of agent-based and learning-based fiscal policy design [39], [40]. The policy vector is $\Theta = \{r_0, \alpha, \beta, \gamma, \delta, \eta, \theta, \lambda, \mu, \nu, \omega, \rho, \kappa, \xi\}$, and the screening problem is the constrained scalarization

$$\Theta^* = \arg \max_{\Theta \in \mathcal{T}} \sum_k \omega_k \tilde{U}_k(\Theta), \quad (19)$$




Fig. 4. Distribution waterfall. Only the attributed AI and robotic shares enter the contribution base; human, data, and conventional capital shares retain their existing fiscal treatment. Shares satisfy $\omega + \kappa + \sigma_T \leq 1$ by Proposition 5.

Algorithm 1 CivicDividendOS annual fiscal cycle

Require: Registry $\{\text{AEAP}_j\}$, episodes $\{\mathcal{P}_i\}$, policy vector Θ

Ensure: Liabilities $\{\text{SAC}_{ij}\}$, fund state F_{t+1} , distributions

- 1: Refresh registry: validate owner, jurisdiction, task class, resources
 - 2: **for all** episodes $p \in \bigcup_i \mathcal{P}_i$ **do**
 - 3: Estimate coalition values $v_p(S)$ for all $S \subseteq \mathcal{F}$
 - 4: Compute Shapley shares $\phi_{p,f}$ by (4); post to ledger
 - 5: Assign mode Z_p and episode scores
 - 6: **end for**
 - 7: **for all** entities i **do**
 - 8: Aggregate $VA_i^{A,R}$; form indices $S_i, A_i^{\text{aug}}, \dots$
 - 9: $\text{ACB}_i \leftarrow VA_i^{A,R} - C_i^{\text{all}}$ ▷ Eq. (6)
 - 10: $r_i \leftarrow \text{clip}(\tilde{r}_i, r_{\min}, r_{\max})$; $\text{SAC}_i \leftarrow r_i \text{ACB}_i$
 - 11: Apportion $\text{SAC}_{ij} \leftarrow \text{Nexus}_{ij} \text{SAC}_i$ ▷ Eq. (16)
 - 12: **end for**
 - 13: $\text{RevSAC} \leftarrow \sum_{i,j} \text{SAC}_{ij}$; verify $\omega + \kappa + \sigma_T \leq 1$
 - 14: Apply shield $\Delta\tau_L$ by (10); disburse transition support
 - 15: $F_{t+1} \leftarrow F_t(1 + r^f) + \omega \text{RevSAC} + E_t - \rho r^f F_t$
 - 16: Distribute $\rho r^f F_t$ across the four channels
 - 17: Twin proposes Θ' ; publish, review, and adopt or reject
-

subject to a revenue floor $\text{Rev}(\Theta) \geq \underline{\text{Rev}}$, an innovation floor $I(\Theta) \geq \underline{I}$, a leakage ceiling $L(\Theta) \leq \bar{L}$, and the boundedness and feasibility constraints of Propositions 3 and 5, where each $\tilde{U}_k$ is a normalized outcome in growth, revenue, welfare, employment, innovation, inequality, and leakage, and the weights ω_k are published rather than inferred.

Two governance constraints on (19) are essential, and both are properties of the institution rather than of the model. First, the twin must report tax incidence explicitly, because a nominal contribution levied on an AI operator may ultimately fall on shareholders, workers, suppliers, or consumers, so policy cannot be evaluated from statutory rates alone. Second, the

twin screens and proposes but does not autonomously set rates: adoption requires publication of Θ , the weights ω_k , and the simulation code; an appeal route for entities contesting their attribution or classification; and periodic external audit of the ledger. Without these constraints an adaptive fiscal system becomes an unaccountable one, and the transparency requirement is what distinguishes the proposal from automated rate setting.

X. WORKED NUMERICAL EXAMPLE

We trace a single entity end to end. An urban logistics company previously employing 100 human dispatchers and drivers introduces AI dispatch agents and autonomous delivery robots. A conventional tax system observes only lower payroll, fewer payroll contributions, higher firm profit, and larger investment in automated capital, and a simple robot tax would charge each robot a fixed annual amount irrespective of whether it replaced a worker, increased safety, or created new demand.

A. Attribution

Let annual value added be $V = 40.0$ monetary units (MU). To make the attribution reproducible rather than asserted, we specify the coalition value function explicitly as a game with pairwise complementarities,

$$v(S) = \sum_{f \in S} a_f + \sum_{\{f, f'\} \subseteq S} b_{ff'}, \quad (20)$$

with singleton weights $a = (a_H, a_A, a_R, a_D, a_K) = (0.220, 0.135, 0.220, 0.020, 0.045)$ and pair weights $b_{HA} = 0.10$, $b_{AD} = 0.05$, $b_{RK} = 0.06$, $b_{AR} = 0.06$, $b_{HR} = 0.04$, $b_{HK} = b_{AK} = 0.02$, $b_{DK} = 0.01$, normalized so that $v(\mathcal{F}) = 1$. The game (20) is monotone, satisfying Assumption 3. For this class the Shapley value admits the closed form $\psi_f = a_f + \frac{1}{2} \sum_{f'} b_{ff'}$, since each singleton term accrues entirely to its factor and each pair term is split equally. Evaluating gives

$$\begin{aligned} \phi_H &= 0.30, & \phi_A &= 0.25, & \phi_R &= 0.30, \\ \phi_D &= 0.05, & \phi_K &= 0.10, \end{aligned} \quad (21)$$

which sum to unity as required by Proposition 1. Summing the AI and robotic entries of (21), the share of attributable value is $VA^{A,R} = 0.55V = 22.0$ MU. With allowable compute, energy, maintenance, and depreciation costs of $C^{\text{all}} = 8.4$ MU, the base is $\text{ACB} = 13.6$ MU.

B. Rate and Liability

Table V lists the calibration. In the substitution-intensive configuration the raw score is $\tilde{r} = 0.116$, which lies inside $[r_{\min}, r_{\max}]$ so $r = 0.116$ and $\text{SAC} = 1.578$ MU. In an augmentation-intensive counterfactual with the same base, in which substitution falls to $S = 0.15$, displacement to $E^{\text{disp}} = 0.10$ and augmentation rises to $A^{\text{aug}} = 0.85$, the raw score falls to $\tilde{r} = 0.0095$ and the liability to 0.129 MU. The same firm with the same automation value added thus faces a liability differing by a factor of 12.2 depending on measured effect

TABLE V

CALIBRATION AND RESULTS FOR THE WORKED EXAMPLE OF SECTION X. INDICES AND WEIGHTS ARE ILLUSTRATIVE POLICY SETTINGS; THE RESULTING LIABILITIES FOLLOW DETERMINISTICALLY FROM (8)–(9).

Quantity	Substitution	Augmentation
<i>Rate weights</i>		
$r_0; r_{\min}; r_{\max}$	0.08; 0.00; 0.25	
$\alpha, \beta, \gamma, \delta, \eta$	0.10, 0.06, 0.05, 0.04, 0.05	
$\theta, \lambda, \mu, \nu$	0.09, 0.07, 0.05, 0.06	
<i>Indices</i>		
S_i	0.62	0.15
C_i^{rent}	0.35	0.35
E_i^{disp}	0.48	0.10
$X_i^{\text{ext}}; R_i^{\text{rev}}$	0.10; 0.55	0.10; 0.55
A_i^{aug}	0.40	0.85
$T_i^{\text{train}}; B_i^{\text{broad}}; N_i^{\text{new}}$	0.55; 0.20; 0.30	0.55; 0.20; 0.30
<i>Outcome</i>		
ACB (MU)	13.60	13.60
$\tilde{r}_i$	0.1160	0.0095
r_i	0.1160	0.0095
SAC _i (MU)	1.578	0.129
Fixed robot tax (MU)	0.240	0.240

on human work, which is precisely the discrimination a fixed robot tax cannot express: a levy of 0.004 MU on each of 60 robots yields 0.240 MU in both configurations. (Annual value added in this example is $V = 40.0$ MU, so the monetary unit is not millions; an earlier version of this example stated the per-robot levy in units inconsistent with V .)

C. Allocation

Revenue is not paid by the robots but by the firm that owns or beneficially uses the automation. With $\omega = 0.40$, $\kappa = 0.25$, $\sigma_T = 0.15$ and a general-revenue residual of 0.20, satisfying the feasibility condition of Proposition 5, the substitution-case liability of 1.578 MU allocates to 0.631 MU into the fund, 0.394 MU to the earned-income shield, 0.237 MU to immediate transition insurance, and 0.316 MU to general revenue. The long-run allocation is therefore that humans receive wages plus a lower labour tax plus a dividend, owners receive profit net of corporate tax and the contribution plus their private return, the public receives the contribution plus public equity returns, and future workers receive the human-capability endowment.

XI. EVALUATION PROTOCOL AND REPRODUCIBILITY

A. Testbed and Baselines

The framework should be evaluated in a heterogeneous-agent macroeconomic and urban-production simulation containing households differing in income, skills, wealth, and automation exposure; firms with heterogeneous productivity and market power; conventional, AI-agent, and robotic capital; government tax and transfer accounts; a public automation wealth fund; and multiple domestic and cross-border sectors. Table VI states the configuration used.

CivicDividendOS should be compared against seven arms: (B0) the existing labour and capital tax system, (B1) a fixed robot tax, (B2) a broad capital-income tax increase, (B3) UBI funded through general taxation, (B4) an automation

TABLE VI

PROPOSED EVALUATION TESTBED CONFIGURATION. THESE ARE THE SETTINGS ACTUALLY USED TO PRODUCE TABLE VII.

Dimension	Setting
Households	10,000; lognormal skill and wealth, exposure negatively correlated with skill
Firms	500; lognormal productivity, beta labour weight
Factors	Labour L and a composite machine factor M
Production	CES in (L, M) , $\sigma = 1.5$, span of control $\gamma = 0.85$
Driver	Machine price falls 0.6% per quarter
Labour supply	Elasticity $\varepsilon_L = 0.30$ to the net-of-tax wage
Jurisdictions	1 (the nexus rule of Sec. VII is not exercised)
Horizon	200 quarters (50 years) plus 40-quarter burn-in
Replications	50 seeds per arm; common random numbers
Attribution	Exact Shapley over $\{L, M\}$, $\psi_L + \psi_M = Y_f$
Deduction share	$\varphi = 1$ (arm B6), $\varphi = 1/2$ (arm B6c)
Budget rule	τ_L clears the budget; spending fixed at 35% of Y
Uncertainty	Paired bootstrap 95% CIs over seeds, 2000 resamples

tax plus targeted retraining, (B5) a sovereign or public wealth fund without factor-origin taxation, and (B6) the integrated framework proposed here, which we run in two variants distinguished only by the cost-deduction share of Proposition 2: B6 with the base exactly as specified and B6c with the corrected base. Every arm finances the same baseline public spending, but the arms are *not* revenue-matched: each also finances its own programme spending, so total spending differs across arms by up to two percentage points of output. Contrasts therefore reflect fiscal stance as well as instrument design, and the revenue-normalised comparison of Section XI-D is the one to read when ranking instruments.

B. Model Specification

The testbed is a discrete-time model of $N_F = 500$ heterogeneous firms and $N_H = 10,000$ heterogeneous households run at quarterly frequency for 40 burn-in and 200 reported quarters. Firm f produces $Y_f = A_f X_f^\gamma$ with $X_f = [\alpha_f L_f^\sigma + (1 - \alpha_f) M_f^\sigma]^{1/\sigma}$, where M denotes combined AI-agent and robotic services, $\sigma = 1/(1 - \varrho) = 1.5$ is the elasticity of substitution between labour and machines, and $\gamma = 0.85$ gives decreasing returns and hence positive pure profit. Productivity A_f and the labour weight α_f are drawn from lognormal and beta distributions respectively, and firms take the wage w and the machine price p_M as given. The machine price falls at 0.6% per quarter, which is the single exogenous driver of the transition.

Households differ in skill, initial wealth, and automation exposure, with exposure negatively correlated with skill. Labour is supplied in efficiency units with elasticity $\varepsilon_L = 0.30$ with respect to the net-of-tax wage, and the wage clears the labour market each quarter. Households whose exposure is high lose efficiency units as automation advances, at rate λ_d per unit of exposure and adoption growth, which is the model's displacement channel; retraining expenditure restores efficiency units in proportion to exposure. Savings rates rise with the wealth rank, which generates a wealth distribution substantially more unequal than the income distribution.

Two features are essential to interpreting the results. First, the model implements the paper's own attribution rule rather

than a proxy for it: ψ_M is computed exactly from (4) for the two-factor game, so $\psi_L + \psi_M = Y_f$ holds identically and the contribution base is the attributed automation value added net of the admitted share φ of machine cost. Second, every arm finances an identical path of baseline public spending, fixed at 35% of output, plus its own programme spending, with the labour income tax adjusting each quarter to clear the government budget. The labour tax rate is therefore an *outcome* rather than a parameter, and the Human Earned-Income Shield of Section VI appears in the results as the reduction in the labour tax that automation revenue makes possible. Arms therefore differ both in the instruments used and in how much revenue those instruments raise, which is why output effects alone do not rank them.

The arms are those of Section XI, with arm B6 split into two variants that differ only in the cost-deduction share of Proposition 2: B6 uses the base exactly as specified in (6) with $\varphi = 1$, and B6c uses the corrected base with $\varphi = 1/2$. Arms B1 and B4 levy a genuine wedge on the machine price, whereas B6 and B6c levy the contribution on the base ACB, whose marginal effect on the machine decision is $r_i \partial \text{ACB} / \partial M$ by Proposition 2. All arms share common random numbers across 50 seeds, so contrasts are paired and confidence intervals are computed by bootstrap over seeds.

C. Metrics

The metrics reported in Table VII are output, hours, the terminal labour share, income and wealth Gini coefficients, the poverty rate, the labour tax rate required to clear the budget, the fund-to-output ratio, and the per-capita dividend. A full study should additionally report: GDP and productivity; AI and robot adoption; employment and hours; wage share and capital share; household disposable income; income and wealth Gini coefficients; poverty rate; revenue volatility; payroll and social-insurance funding; investment and innovation; firm entry and exit; public-fund value; citizen dividend per capita; worker-transition duration; human-capability investment; tax avoidance and leakage; and intergenerational wealth distribution. Sensitivity analysis should vary AI and robot costs, substitution elasticities, market concentration, international capital mobility, tax compliance, and the speed of technological progress. Effects should be reported with uncertainty, in the form $\Delta \text{Gini} = -0.044$, $\text{CI}_{95\%} = [-0.058, -0.031]$, rather than as point estimates.

D. Results

Table VII reports the campaign and Table VIII the paired contrasts. All numerical entries in these tables are generated directly from the simulation output by `scripts/make_tables.py` and included with `\input`; none is transcribed by hand.

Two framing points come first, because they govern how everything below should be read.

The confidence intervals are replication intervals. Arms share common random numbers, so contrasts are paired and the bootstrap intervals over seeds are extremely tight. That is a statement about the reproducibility of the simulation at one parameter vector, and not about the certainty of any policy

conclusion. The sensitivity analysis of Section XI-E moves the central estimate by tens of percentage points, which no replication interval here can see.

The arms are not revenue-matched. Each finances the same baseline public spending plus its own programme, so total spending differs across arms. Ranking instruments by output effect alone is therefore not valid, and Table IX gives the revenue-normalised comparison.

Six findings stand out, and they do not uniformly favour the proposed framework.

First, the fixed robot tax of arm B1 is by a wide margin the most costly instrument in output terms, at -6.12% relative to baseline, while raising 1.91% of output in revenue. It does support the labour share, which ends at 0.341 against 0.298 in the baseline, but it achieves this by taxing an input at a uniform rate irrespective of whether the automation in question displaced anyone. This is the quantitative counterpart of the attribution gap of Section III.

Second, the framework as literally specified in (6) is close to inert. Arm B6 raises only 0.24% of output and reduces the income Gini by 0.0022 , which is statistically clear and economically negligible. Its output effect is positive, at $+2.45\%$, *precisely because* the base is marginally subsidising in the sense of Proposition 2. A positive output number here is evidence of the design defect, not of merit.

Third, the corrected base repairs the revenue problem but does not deliver dominance. Arm B6c raises 0.88% of output, lowers the required labour tax from 0.381 to 0.370 ($\Delta \tau_L = -0.0111$, $\text{CI}_{95\%} = [-0.0118, -0.0103]$), reduces the income Gini by 0.0072 ($\text{CI}_{95\%} = [-0.0077, -0.0068]$), and raises output by 3.18% and labour demand in efficiency units by 6.78% , holding the terminal labour share at 0.316 against 0.298 . These figures are produced under corrected budget accounting, in which the full contribution liability $r_i \text{ACB}_i$ is debited from firm profit while the marginal decision price remains $p_M + r_i \partial \text{ACB} / \partial M$ as Proposition 2 requires. Under the earlier accounting, in which only the marginal wedge was debited, the same arm showed $+3.30\%$; the difference is the mechanical effect of charging the inframarginal liability to somebody.

Fourth, and decisively against the strong reading of the proposal, B6c is *not* the most efficient instrument once revenue is held constant. Table IX reports revenue raised per unit of deadweight loss on the labour margin. B6c returns 0.45 , against 1.47 for the fixed robot tax of B1 and 0.89 for the automation levy of B4. B6c looks good on output largely because it raises little revenue. Furthermore, arm B9, which finances the same programme with a lump-sum levy rather than with the attributed contribution, attains $+3.03\%$ output against B6c's $+3.18\%$. Most of the output advantage is therefore a fiscal-stance effect rather than a return to the attribution machinery, which is the central claim the framework needs and does not establish here.

Fifth, targeted transfers remain the better instrument for near-term distributional objectives. Arm B3 attains the lowest income Gini at 0.3147 and much the lowest poverty rate at 9.95% , outcomes that B6c does not approach, and it does so immediately rather than over decades. The cost is the highest labour tax of any arm at 0.413 . A fund whose half-life is measured in decades cannot substitute for transfers within a

political cycle: the dividend in B6c reaches only 0.26% of output after fifty years.

Sixth, arm B4, a flat automation levy recycled into retraining, delivers the largest gain in labour demand of any arm, at +8.52% in efficiency units, for an output cost of only -0.72%. Much of the benefit attributed to the attribution machinery can be obtained by a simpler instrument coupled to an active labour-market policy.

A note on the employment column. The quantity reported here is aggregate labour demand measured in *efficiency units*, $\sum_f L_f$. It is not hours and it is not employment. A worker whose efficiency has been reduced by displacement contributes less to it while remaining fully employed, and retraining expenditure raises it directly. The model has no extensive employment margin — labour supply is a smooth function of the net-of-tax wage and no agent is ever unemployed — so it cannot speak to employment in the ordinary sense. Table X reports efficiency units, a headcount above the displacement floor, and the wage separately.

E. Sensitivity and the Region of Validity

The results of Section XI-D are computed at a single point, $\sigma = 1.5$ and $\gamma = 0.85$. Neither parameter is calibrated, and the Metrics discussion of Section XI requires that sensitivity to them be reported. We therefore sweep both.

Table XII reports arm B6c's output effect across $\sigma \in \{1.2, 1.5, 2.0, 3.0, 5.0, 8.0\}$ and $\gamma \in \{0.75, 0.85, 0.95\}$, at ten seeds per cell under common random numbers.

The output advantage is positive in only four of the eighteen cells. At $\gamma = 0.85$ it is +4.78% at $\sigma = 1.2$ and +3.01% at $\sigma = 1.5$, but -2.23% at $\sigma = 2.0$ and -24.99% at $\sigma = 8.0$. The sign therefore reverses between $\sigma = 1.5$ and $\sigma = 2.0$, and the region in which the advantage holds is approximately $\sigma \lesssim 1.7$ at $\gamma = 0.85$. Since the empirical literature places the labour-capital elasticity of substitution across a range that comfortably spans this boundary, **the result must be read as conditional on a low elasticity and not as a general property of the instrument.** At $\gamma = 0.95$, close to constant returns, firm scale becomes explosive and the model is better regarded as outside its domain of validity than as delivering a policy result; that row is reported as a boundary diagnostic.

The same sweep clarifies Proposition 2. Arm B6, with full deduction, shows a small positive output effect almost everywhere in the grid, which is exactly what a marginally subsidising base predicts. The proposition is therefore supported by the sweep, while the proposal's headline advantage is not.

1) *Two further reversals:* The elasticity of substitution is not the only dimension along which the sign turns. Two others matter as much, and neither was reported in the earlier version of this work.

Horizon. The reporting window of 200 quarters is a choice, not a result. Arm B6c's output effect is +1.78% at 80 quarters, +2.81% at 200, -2.23% at 320 and -11.20% at 400. The advantage is therefore transitional: it accrues while the fund is accumulating and the machine price has not yet fallen far, and it reverses once the contribution has been levied on a large

automated base for long enough. Reporting a fifty-year window and stopping there flatters the instrument.

Speed of automation. The machine price decline of 0.6% per quarter is assumed rather than calibrated, and it drives the entire transition. At 0.2% the output effect is +2.24% and at 0.4% it is +2.85%, but at 1.0% it is -7.34% and at 1.5% it is -22.45%. The instrument performs well against slow automation and badly against fast automation, which is close to the opposite of the case usually made for it.

Taken with the elasticity sweep, the honest summary is that arm B6c's measured advantage requires a low elasticity of substitution, a moderate pace of automation, and a horizon of roughly fifty years. It is a statement about a region of the parameter space, and the region does not obviously contain the world.

F. Cross-Border Leakage and Classification Error

The single-jurisdiction testbed of the earlier version could not speak to leakage at all. Enabling the three-jurisdiction nexus of Section VII, with rate multipliers (1.0, 0.5, 1.5) and a base that relocates towards the cheaper jurisdiction, gives revenue losses of 2.1%, 4.1% and 8.2% at shifting elasticities of 0.2, 0.4 and 0.8. Apportionment remains exhaustive throughout, as Proposition 8 requires: the loss appears as revenue accruing to another jurisdiction, not as a violated identity. Leakage is therefore bounded and roughly linear in the elasticity here, but it is not eliminated, and the magnitude depends entirely on how mobile the four nexus components are in practice — which this model assumes rather than measures.

Classification error behaves as Proposition 4 predicts. At injected error rates of 5%, 10% and 20% the induced revenue error is 3.2%, 6.7% and 13.7% of the clean revenue, against a bound of $\varepsilon(\alpha + \theta)$ that permits 11.9%, 23.8% and 47.5%. The bound holds with room to spare, and the error grows linearly rather than catastrophically. This is the one robustness claim in the paper that survives the corrected testbed intact.

G. Fund Stability in the Testbed

Proposition 6 requires $(1 - \rho)r^f < g$ for the fund-to-output ratio to converge. The testbed sets $\rho = 0.60$ and $r^f = 0.05$, giving $(1 - \rho)r^f = 0.0200$, while the model's own realised output growth across fifty seeds ranges from 0.0147 to 0.0336 with a mean of 0.0231.

The condition therefore fails on a substantial minority of runs: 13 of 50 seeds are on a divergent path. On those runs the fund-to-output ratio has no finite limit, and the terminal level reported in Table VII is a point on a trajectory rather than a steady state. It should not be compared with the analytic f^* of (14), and the 71-year half-life of Table IV assumes $g = 0.03$, which this testbed does not deliver. Table IV's arithmetic is itself correct; the difficulty is the growth rate at which it is evaluated.

We report this rather than reparameterising to remove it. Selecting ρ or r^f so that the condition holds would make the fund look convergent by construction and would leave the reader no way to tell how close to the boundary the design sits. The honest summary is that the fund as calibrated is on

TABLE VII

SIMULATION RESULTS: MEANS OVER 50 SEEDS UNDER COMMON RANDOM NUMBERS, AVERAGED ACROSS THE 200 REPORTED QUARTERS EXCEPT THE WAGE SHARE AND THE FUND (TERMINAL VALUES). OUTPUT AND LABOUR DEMAND ARE PAIRED PERCENTAGE DIFFERENCES AGAINST ARM B0. *Labour demand is measured in efficiency units and is neither hours nor employment.* THE LABOUR TAX RATE CLEARS THE GOVERNMENT BUDGET AND IS AN OUTCOME. PRODUCED UNDER CORRECTED BUDGET ACCOUNTING. THESE VALUES ARE **SYNTHETIC** OUTPUTS OF THE MODEL OF SECTION XI-B, NOT EMPIRICAL ESTIMATES.

Arm	Policy regime	Δ GDP (%)	Labour (eff., %)	Wage share	Gini (inc.)	Gini (wealth)	τ_L req.	Fund (%GDP)	Div. (%GDP)
B0	Existing labour/capital tax system	+0.00	+0.00	0.298	0.3296	0.5575	0.381	0.0	0.00
B1	Fixed robot tax	-6.12	+3.34	0.341	0.3245	0.5558	0.325	0.0	0.00
B2	Broad capital-income tax increase	+0.84	+1.20	0.298	0.3255	0.5535	0.356	0.0	0.00
B3	UBI via general taxation	-1.05	-1.55	0.297	0.3147	0.5569	0.413	0.0	0.00
B4	Automation tax plus retraining	-0.72	+8.52	0.338	0.3247	0.5556	0.352	0.0	0.00
B5	Public wealth fund, no factor-origin tax	-0.10	-0.14	0.298	0.3283	0.5575	0.384	7.4	0.17
B6	CivicDividendOS, base as specified	+2.45	+1.56	0.297	0.3274	0.5569	0.385	4.3	0.08
B6c	CivicDividendOS, corrected base	+3.18	+6.78	0.316	0.3224	0.5558	0.370	15.1	0.26
B7	Optimal linear capital tax	+0.41	+0.60	0.298	0.3275	0.5555	0.369	0.0	0.00
B8	Consumption-tax shift	+1.01	+1.45	0.299	0.3282	0.5562	0.350	0.0	0.00
B9	Lump-sum-financed CivicDividendOS	+3.03	+2.38	0.298	0.3351	0.5573	0.368	4.3	0.08
B10	EITC-style wage subsidy	-0.23	-0.35	0.298	0.3213	0.5569	0.389	0.0	0.00
B11	CivicDividendOS, no cost deduction ($\phi = 0$)	-3.13	+2.07	0.319	0.3255	0.5559	0.345	0.0	0.00
B12	Thuemmel-style optimal robot tax	-2.54	+1.34	0.313	0.3275	0.5568	0.359	0.0	0.00

TABLE VIII

PAIRED DIFFERENCES AGAINST ARM B0 WITH BOOTSTRAP 95% CONFIDENCE INTERVALS OVER 50 SEEDS (2000 RESAMPLES). THESE ARE *replication* INTERVALS: THEY QUANTIFY MONTE-CARLO VARIATION AT ONE FIXED PARAMETER VECTOR AND SAY NOTHING ABOUT PARAMETER UNCERTAINTY. SEE SECTION XI-E.

Arm	Δ Gini	95% CI	$\Delta\tau_L$	95% CI
B1	-0.0051	[-0.0056, -0.0046]	-0.0558	[-0.0577, -0.0537]
B2	-0.0042	[-0.0044, -0.0039]	-0.0257	[-0.0262, -0.0252]
B3	-0.0149	[-0.0154, -0.0145]	+0.0322	[+0.0315, +0.0329]
B4	-0.0050	[-0.0053, -0.0046]	-0.0288	[-0.0296, -0.0278]
B5	-0.0013	[-0.0013, -0.0013]	+0.0029	[+0.0028, +0.0030]
B6	-0.0022	[-0.0023, -0.0021]	+0.0033	[+0.0032, +0.0034]
B6c	-0.0072	[-0.0077, -0.0068]	-0.0111	[-0.0118, -0.0103]
B7	-0.0022	[-0.0023, -0.0021]	-0.0127	[-0.0129, -0.0125]
B8	-0.0014	[-0.0015, -0.0014]	-0.0310	[-0.0315, -0.0304]
B9	+0.0055	[+0.0053, +0.0056]	-0.0137	[-0.0140, -0.0134]
B10	-0.0083	[-0.0084, -0.0083]	+0.0075	[+0.0073, +0.0077]
B11	-0.0041	[-0.0046, -0.0037]	-0.0360	[-0.0377, -0.0343]
B12	-0.0022	[-0.0023, -0.0020]	-0.0220	[-0.0225, -0.0214]

TABLE IX

REVENUE-NORMALISED COMPARISON. RANKING INSTRUMENTS BY OUTPUT EFFECT ALONE IS NOT VALID WHEN THEY RAISE MATERIALLY DIFFERENT REVENUE: AN INSTRUMENT THAT RAISES ALMOST NOTHING WILL LOOK EFFICIENT FOR NO BETTER REASON THAN THAT IT DOES ALMOST NOTHING. DWL IS THE HARBERGER TRIANGLE ON THE LABOUR MARGIN.

Arm	Δ GDP (%)	Revenue (%Y)	DWL (%Y)	Rev/DWL
B0	+0.00	0.000	2.307	0.000
B1	-6.12	1.907	1.488	1.472
B2	+0.84	0.000	1.954	0.000
B3	-1.05	0.000	2.789	0.000
B4	-0.72	1.476	1.924	0.892
B5	-0.10	0.000	2.355	0.000
B6	+2.45	0.238	2.352	0.109
B6c	+3.18	0.877	2.186	0.451
B7	+0.41	0.000	2.129	0.000
B8	+1.01	0.000	1.885	0.000
B9	+3.03	0.239	2.113	0.123
B10	-0.23	0.000	2.404	0.000
B11	-3.13	1.486	1.783	0.941
B12	-2.54	0.800	1.974	0.472

TABLE X

LABOUR-MARKET QUANTITIES, SEPARATED. EFFICIENCY UNITS ARE AGGREGATE LABOUR DEMAND $\sum_f L_f$; HEADCOUNT COUNTS WORKERS ABOVE THE DISPLACEMENT FLOOR. THE MODEL HAS NO EXTENSIVE EMPLOYMENT MARGIN — LABOUR SUPPLY IS A SMOOTH FUNCTION OF THE NET-OF-TAX WAGE AND NO AGENT IS EVER UNEMPLOYED — SO NEITHER COLUMN IS AN UNEMPLOYMENT RATE, AND THE QUANTITY REPORTED AS “HOURS” IN AN EARLIER VERSION OF THIS TABLE WAS THE FIRST OF THESE.

Arm	Eff. units	Δ (%)	Headcount	Δ (%)	Wage
B0	6470.0	+0.00	10000.0	+0.00	0.4980
B1	6685.2	+3.34	10000.0	+0.00	0.4842
B2	6547.9	+1.20	10000.0	+0.00	0.4968
B3	6369.9	-1.55	10000.0	+0.00	0.4995
B4	7020.0	+8.52	10000.0	+0.00	0.4816
B5	6460.8	-0.14	10000.0	+0.00	0.4981
B6	6570.8	+1.56	10000.0	+0.00	0.4979
B6c	6907.9	+6.78	10000.0	+0.00	0.4897
B7	6508.7	+0.60	10000.0	+0.00	0.4974
B8	6563.7	+1.45	10000.0	+0.00	0.4965
B9	6624.2	+2.38	10000.0	+0.00	0.4971
B10	6447.4	-0.35	10000.0	+0.00	0.4983
B11	6603.5	+2.07	10000.0	+0.00	0.4920
B12	6556.3	+1.34	10000.0	+0.00	0.4924

TABLE XI

ECONOMIC INCIDENCE OF THE AUTOMATION LEVY, APPORTIONED BY FACTOR INCOME SHARES. STATUTORY INCIDENCE IS ON THE FIRM IN EVERY CASE. REPORTING THIS SPLIT IS A GOVERNANCE REQUIREMENT OF SECTION IX, BECAUSE A CONTRIBUTION LEVIED ON AN OPERATOR MAY ULTIMATELY FALL ON SHAREHOLDERS, WORKERS, SUPPLIERS OR CONSUMERS.

Arm	Borne by labour (%)	Borne by capital (%)	Revenue (%Y)
B1	76.3	23.7	1.907
B4	76.1	23.9	1.476
B6	76.6	23.4	0.238
B6c	76.7	23.3	0.877
B9	76.6	23.4	0.239
B11	77.0	23.0	1.486
B12	75.2	24.8	0.800

XII. DISCUSSION

A. Incidence

The statutory taxpayer is the owner of record, but statutory and economic incidence differ. In the short run, with automation capital fixed, the contribution falls largely on quasi-rents and

the edge of stability, and that a design intended to run for a century should not be.

TABLE XII

SENSITIVITY OF THE OUTPUT EFFECT TO THE ELASTICITY OF SUBSTITUTION σ AND THE SPAN OF CONTROL γ , TEN SEEDS PER CELL UNDER COMMON RANDOM NUMBERS, WITH $\rho = 1 - 1/\sigma$. THE $\gamma > \rho$ COLUMN RECORDS WHETHER THE MARGINAL-INCIDENCE INEQUALITY OF PROPOSITION 2 HOLDS IDENTICALLY. **ARM B6C'S ADVANTAGE IS POSITIVE IN ONLY FOUR OF THE EIGHTEEN CELLS, AND THE SIGN REVERSES BETWEEN $\sigma = 1.5$ AND $\sigma = 2.0$ AT $\gamma = 0.85$. THE $\gamma = 0.95$ ROW IS A BOUNDARY DIAGNOSTIC: NEAR CONSTANT RETURNS FIRM SCALE BECOMES EXPLOSIVE AND THE MODEL IS OUTSIDE ITS DOMAIN OF VALIDITY.**

σ	γ	ρ	$\gamma > \rho$	B6 (%)	B6c (%)	B6c 95% CI
1.2	0.75	0.167	yes	+3.37	+3.29	[+3.21, +3.38]
1.2	0.85	0.167	yes	+3.44	+4.56	[+4.23, +4.82]
1.2	0.95	0.167	yes	+5.99	-0.36	[-3.53, +2.26]
1.5	0.75	0.333	yes	+2.01	+2.32	[+2.24, +2.40]
1.5	0.85	0.333	yes	+2.40	+2.81	[+1.83, +3.48]
1.5	0.95	0.333	yes	+2.22	-48.87	[-55.75, -41.92]
2.0	0.75	0.500	yes	+1.37	+1.15	[+1.08, +1.20]
2.0	0.85	0.500	yes	+2.11	-2.24	[-4.79, -0.27]
2.0	0.95	0.500	yes	+0.04	-61.34	[-63.92, -58.76]
3.0	0.75	0.667	yes	+1.09	-0.03	[-0.11, +0.05]
3.0	0.85	0.667	yes	+1.69	-11.66	[-15.44, -8.42]
3.0	0.95	0.667	yes	+0.00	-61.60	[-64.12, -58.85]
5.0	0.75	0.800	no	+0.93	-0.91	[-1.04, -0.79]
5.0	0.85	0.800	yes	+1.17	-20.45	[-23.99, -17.36]
5.0	0.95	0.800	yes	+0.00	-61.70	[-64.22, -58.93]
8.0	0.75	0.875	no	+0.90	-1.36	[-1.50, -1.22]
8.0	0.85	0.875	no	+0.94	-24.79	[-28.00, -21.98]
8.0	0.95	0.875	yes	+0.00	-61.75	[-64.26, -58.98]

therefore on owners. As capital becomes mobile over the policy horizon, part of the burden shifts to immobile factors, which in an open economy means labour and land. This is the standard argument against source-based capital taxation, and it applies here with the important qualification that the base is attributed automation value added rather than total capital income, so the distortion is concentrated on the margin between automating and not automating a given task rather than on capital accumulation generally. The Diamond-Mirrlees production-efficiency result implies that such an input tax is justified only to the extent it corrects an externality [20]; the framework's rationale is accordingly the uninternalized transition cost documented in [16], not revenue as such.

B. Avoidance and Gaming

Four avoidance channels deserve attention. Strategic task decomposition can shift value across episode boundaries to reduce measured substitution, which argues for aggregation at the entity rather than episode level when computing indices. Relabelling substitution as augmentation is bounded in fiscal consequence by Proposition 4 but not eliminated, and requires audit of the employment panels underlying S_i . Cost inflation in C^{all} directly reduces the base and needs the same verification machinery as existing depreciation rules. Offshoring of deployment is addressed by the nexus of (16) only to the extent that usage, revenue, users, and physical operations are themselves verifiable.

C. Deployment Considerations

A realistic path is phased. A registration-only phase establishes the passport and ledger with no liability attached, which builds the measurement base and permits the classifier to be

validated before it carries fiscal consequence. A disclosure phase publishes attributed shares and shadow liabilities. Only a third phase attaches rates, initially at a low r_{max} , with the fund capitalized first from equity contributions under Section VI rather than cash. This sequencing matters because the framework's measurement requirements are its binding constraint, and attaching liability to an unvalidated ledger would import estimation error directly into legal obligations.

XIII. LIMITATIONS AND THREATS TO VALIDITY

Several limitations qualify the contribution. First, Assumption 2 is strong: coalition values $v_p(S)$ are counterfactual and not directly observed, so attribution inherits the error of the estimator, and by the linearity of (4) an error of ϵ_v in coalition values propagates to at most ϵ_v in each attributed share. Validating these counterfactuals is the single largest empirical obstacle to implementation. Second, and most importantly, the results of Section XI-D are outputs of the model of Section XI-B, not empirical estimates. That model is deliberately stylized: it has one composite machine factor rather than the five factors the framework distinguishes, no explicit data or intangible input, a single good, no entry or exit, no international sector despite the nexus rule of Section VII, and an exogenous machine price rather than endogenous innovation. Its magnitudes should be read as internal comparisons across arms under a common structure, and the sign and ordering of effects are more credible than their size. Calibration against national accounts and robot-stock data remains to be done. Third, the analytical results of Section VI rely on Assumption 6, under which returns and growth are constant; realistic return volatility would make the dividend procyclical unless the payout rule is smoothed, which is a design question we do not resolve. Fourth, the measurement burden falls hardest on small firms, so the de minimis threshold of Assumption 1 is doing substantial work and its calibration determines whether the system is administrable. Fifth, the nexus rule requires international coordination that may not be forthcoming, and unilateral adoption would invite the same leakage the rule is designed to prevent [42]; the testbed has a single jurisdiction and so cannot speak to leakage at all. Sixth, task-level metering raises privacy concerns where episodes are attributable to identifiable workers, and the framework as specified does not include the disclosure limits such metering would require. Finally, the political economy of an adaptive rate is unaddressed: a rate that responds to measured substitution creates an incentive to contest measurement, and the appeal machinery of Section IX would carry more weight than the estimator itself.

XIV. CONCLUSION

The AI and robotics era creates a distribution problem before it creates a purely technical tax problem. If increasingly autonomous productive systems generate a rising share of economic value while wages remain the principal source of household income and payroll taxation remains central to public finance, productivity can rise while purchasing power, social-insurance finance, and wealth distribution become more fragile.

CivicDividendOS addresses this mismatch by separating three questions that are frequently conflated: who created economic value, who is legally responsible for tax, and who should share in the returns. The Factor-Origin Ledger answers the first with a Shapley-consistent allocation that is budget balanced by construction. The Autonomous Economic Activity Passport answers the second by binding machine activity to a human or corporate owner of record without conferring legal personality. The bounded adaptive contribution, earned-income shield, wealth fund, and four-channel waterfall answer the third, with an explicit cap on the marginal burden, a bounded sensitivity to classifier error, and a closed-form stability condition for the fund.

Two results are of independent interest. The fiscal cost of misclassification is Lipschitz in the classifier's error rate, so a noisy measurement layer bounds rather than propagates uncertainty, and the weights can be chosen against an explicit tolerance. The steady-state per-capita dividend is decreasing in the fund's payout ratio whenever the fund return exceeds output growth, so the political pressure to raise payouts is precisely the pressure that erodes the long-run dividend; combined with a convergence half-life measured in decades, this implies that such a fund is an intergenerational instrument rather than a short-run income policy.

The framework does not seek to make robots tax citizens. It creates a measurable and adaptive fiscal bridge between an economy in which humans earn most productive income and one in which a larger share of value may be generated by autonomous capital. Two qualifications follow from the evaluation and should travel with the proposal. The contribution base must deduct costs at no more than half their value, or the instrument subsidises the activity it is meant to price; and a wealth fund is an instrument for the next generation rather than for the current one, so it complements transfers and retraining rather than replacing them. The immediate priority for future work is empirical: validating the coalition-value estimator against observed counterfactuals, extending the testbed to the five-factor, multi-jurisdiction setting the framework assumes, and calibrating it against national accounts so that the internal comparisons reported here can be given empirical magnitudes.

REFERENCES

- [1] D. Acemoglu and P. Restrepo, "Automation and new tasks: How technology displaces and reinstates labor," *Journal of Economic Perspectives*, vol. 33, no. 2, pp. 3–30, 2019.
- [2] E. Brynjolfsson, T. Mitchell, and D. Rock, "What can machines learn and what does it mean for occupations and the economy?" *AEA Papers and Proceedings*, vol. 108, pp. 43–47, 2018.
- [3] T. Eloundou, S. Manning, P. Mishkin, and D. Rock, "GPTs are GPTs: An early look at the labor market impact potential of large language models," *arXiv preprint arXiv:2303.10130*, 2023.
- [4] International Monetary Fund, "Fiscal policy can help broaden the gains of AI to humanity," IMF Blog, Jun. 2024.
- [5] A. Korinek and L. M. Lockwood, "Public finance in the age of AI: A primer," National Bureau of Economic Research, NBER Working Paper 34873, 2026.
- [6] C. Dimitropoulou, "Current income (profit) taxation of artificial intelligence," in *Research Handbook on the Taxation of Robots and Artificial Intelligence*. Edward Elgar Publishing, [UNVERIFIED] Chapter authorship could not be confirmed. The Elgar volume is "Taxing Artificial Intelligence" (2025). If the intended source is Dimitropoulou's monograph, cite instead: C. Dimitropoulou, Robot Taxation: A Normative Tax Policy Analysis, IBFD Doctoral Series 70, 2024.
- [7] D. H. Autor, F. Levy, and R. J. Murnane, "The skill content of recent technological change: An empirical exploration," *The Quarterly Journal of Economics*, vol. 118, no. 4, pp. 1279–1333, 2003.
- [8] D. Acemoglu and P. Restrepo, "Robots and jobs: Evidence from U.S. labor markets," *Journal of Political Economy*, vol. 128, no. 6, pp. 2188–2244, 2020.
- [9] W. Dauth, S. Findeisen, J. Südekum, and N. Woessner, "The adjustment of labor markets to robots," *Journal of the European Economic Association*, vol. 19, no. 6, pp. 3104–3153, 2021.
- [10] C. B. Frey and M. A. Osborne, "The future of employment: How susceptible are jobs to computerisation?" *Technological Forecasting and Social Change*, vol. 114, pp. 254–280, 2017.
- [11] International Monetary Fund, "AI will transform the global economy. let's make sure it benefits humanity," IMF Blog, Jan. 2024.
- [12] OECD, *Artificial Intelligence and the Labour Market in Korea*. Paris: OECD Publishing, 2025.
- [13] J. Guerreiro, S. Rebelo, and P. Teles, "Should robots be taxed?" *The Review of Economic Studies*, vol. 89, no. 1, pp. 279–311, 2022.
- [14] U. Thummel, "Optimal taxation of robots," *Journal of the European Economic Association*, vol. 21, no. 3, pp. 1154–1190, 2023.
- [15] D. Acemoglu, A. Manera, and P. Restrepo, "Does the U.S. tax code favor automation?" National Bureau of Economic Research, NBER Working Paper 27052, 2020, also published in Brookings Papers on Economic Activity, 2020(1), pp. 231–300.
- [16] M. Beraja and N. Zorzi, "Inefficient automation," *The Review of Economic Studies*, vol. 92, no. 1, 2025.
- [17] R. Merola, "Inclusive growth in the era of automation and AI: How can taxation help?" *Frontiers in Artificial Intelligence*, vol. 5, 2022.
- [18] P. Zhang, "Automation, wage inequality and implications of a robot tax," *International Review of Economics & Finance*, vol. 59, pp. 500–509, 2019.
- [19] J. A. Mirrlees, "An exploration in the theory of optimum income taxation," *The Review of Economic Studies*, vol. 38, no. 2, pp. 175–208, 1971.
- [20] P. A. Diamond and J. A. Mirrlees, "Optimal taxation and public production I: Production efficiency," *The American Economic Review*, vol. 61, no. 1, pp. 8–27, 1971.
- [21] A. B. Atkinson and J. E. Stiglitz, "The design of tax structure: Direct versus indirect taxation," *Journal of Public Economics*, vol. 6, no. 1–2, pp. 55–75, 1976.
- [22] K. Hötte, A. Theodorakopoulos, and P. Kourtroumpis, "Automation and taxation," *Oxford Economic Papers*, vol. 76, no. 4, pp. 945–973, 2024.
- [23] International Monetary Fund, "Global economic and financial implications of artificial intelligence," International Monetary Fund, Tech. Rep., 2026, [INCOMPLETE] Confirm series and report number.
- [24] J. Faivre and S. H. Cen, "Taxing artificial intelligence," *arXiv preprint arXiv:2607.02144*, 2026.
- [25] P. Van Parijs and Y. Vanderborght, *Basic Income: A Radical Proposal for a Free Society and a Sane Economy*. Cambridge, MA: Harvard University Press, 2017.
- [26] I. Ortiz, C. Behrendt, A. Acuña-Ulate, and Q. A. Nguyen, "Universal basic income proposals in light of ILO standards: Key issues and global costing," International Labour Organization, Geneva, ESS Working Paper 62, 2018.
- [27] J. Le Grand, "A springboard for new citizens: Universal basic capital," *LSE Public Policy Review*, vol. 1, no. 2, 2020.
- [28] G. Corneo, "Progressive sovereign wealth funds," *Journal of Government and Economics*, vol. 5, p. 100033, 2022.
- [29] T. Piketty, *Capital in the Twenty-First Century*. Cambridge, MA: Harvard University Press, 2014.
- [30] International Forum of Sovereign Wealth Funds, "Alaska permanent fund," Institutional profile, [INCOMPLETE] Access date and URL required.
- [31] Brookings Institution, "The future of tax policy: A public finance framework for the age of AI," Brookings Institution report, Jan. 2026.
- [32] —, "AI growth acceleration versus distributional fairness," Brookings Institution report, May 2026.
- [33] N. Vincent, Y. Li, R. Zha, and B. Hecht, "Mapping the potential and pitfalls of data dividends as a means of sharing the profits of artificial intelligence," *arXiv preprint arXiv:1912.00757*, 2019.
- [34] E. Bax, "Computing a data dividend," *arXiv preprint arXiv:1905.01805*, 2019.
- [35] L. S. Shapley, "A value for n -person games," in *Contributions to the Theory of Games, Volume II*, H. W. Kuhn and A. W. Tucker, Eds. Princeton, NJ: Princeton University Press, 1953, pp. 307–317.
- [36] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

- [37] A. Ghorbani and J. Zou, “Data Shapley: Equitable valuation of data for machine learning,” in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019.
- [38] R. Jia, D. Dao, B. Wang, F. A. Hubis, N. Hynes, N. M. Gürel, B. Li, C. Zhang, D. Song, and C. J. Spanos, “Towards efficient data valuation based on the Shapley value,” in *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2019.
- [39] J. Wang, X. Fang, L. Huang, and Y. Huang, “TaxAgent: How large language model designs fiscal policy,” *arXiv preprint arXiv:2506.02838*, 2025.
- [40] S. Zheng, A. Trott, S. Srinivasa, D. C. Parkes, and R. Socher, “The AI economist: Taxation policy design via two-level deep reinforcement learning,” *Science Advances*, vol. 8, no. 18, 2022.
- [41] P. Fratrič, N. Holzenberger, and D. Restrepo Amariles, “Can AI expose tax loopholes? towards a new generation of legal policy assistants,” *arXiv preprint arXiv:2503.17339*, 2025.
- [42] OECD/G20 Inclusive Framework on BEPS, “Statement on a two-pillar solution to address the tax challenges arising from the digitalisation of the economy,” OECD, Paris, Tech. Rep., 2021.